



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

# A Quick guide to apply Clustering and Tracking to your traces

Presentation slides for MIRI-PPTM. C. Gómez

Barcelona, 29 abril 2015



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

# CLUSTERING

# Outline

## « Clustering

- « Setting up the environment

- « Run

- « Result output files

- « General tips and common issues

## « Tracking

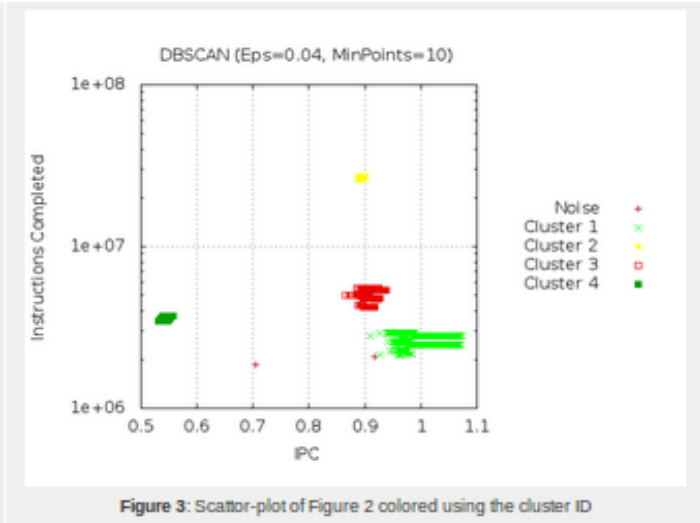
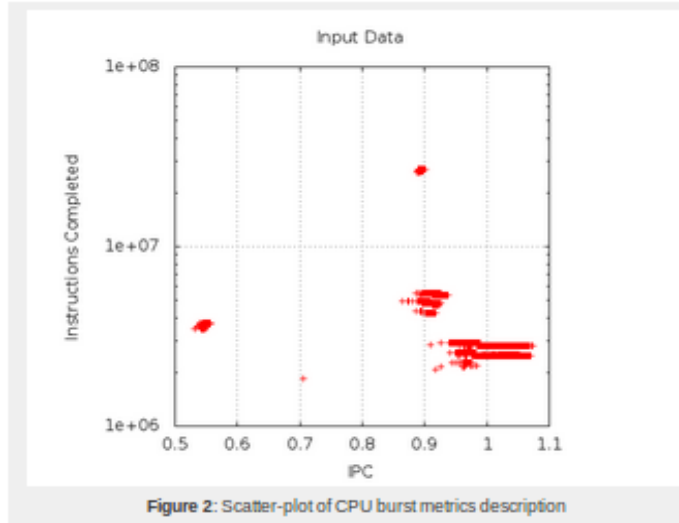
- « Set up and run

- « Tracking GUI




# Clustering

Characterizes computational bursts that are similar and groups them in clusters



Allows you to study the behavior of the clusters separately, identify patterns etc.



# Clustering

« Manual and more details about the tool

[http://www.bsc.es/sites/default/files/public/computer\\_science/performance\\_tools/clustersuite\\_manual.pdf](http://www.bsc.es/sites/default/files/public/computer_science/performance_tools/clustersuite_manual.pdf)

note: If you never heard about or seen any clustering example, you should read the main page info of the tool first.

# Clustering – Setting up the environment

- ⌘ It's recommended to work remotely connected to MN3

- ⌘ Everything is configured automatically with modules

- ⌘ Tracking has no distributable package yet

- ⌘ Log in with: `ssh -X -l <user> mn3.bsc.es`

- ⌘ Load modules: `module load INTEL BSCTOOLS`

- ⌘ Your traces have at least this PAPI counters info

- ⌘ Configure in `extrae.xml`: `PAPI_TOT_INS` , `PAPI_TOT_CYC`

# Clustering – Setting up the environment

## ⌘ Locate the cluster.xml configuration file

⌘ `/apps/CEPBATTOOLS/ClusteringSuite/latest/share/example/cluster.xml`

⌘ Copy it to your working folder, you ~~may~~ **need** to modify it

## ⌘ cluster.xml parameters (try default first!)

⌘ Epsilon – 'radius' in which a burst will be considered similar

⌘ Min points – minimum number of similar points to form a cluster

⌘ Range min/max: Allows to filter very small bursts or focus in a desired area.

# Clustering - Run

- « Applying clustering to your trace

- « Open Paraver

- « Click on the cog icon

- « Application: clustering

- « Select your trace file...

- « Select your cluster.xml file...

- « Algorithm: XML defined (should be DBScan)

- « Press Run!



# Clustering – Output files generated

- ⌘ The process will generate several output files

- ⌘ A New trace `mytrace.clustered.prv`

- ⌘ In paraver use `clustering/clusterID_window.cfg`

- ⌘ A Instructions / IPC plot (this is the cool one!)

- ⌘ Open with `gnuplot mytrace.clustered.IPC.PAPI_TOT_INS.gnuplot`

- ⌘ Clusters duration information

- ⌘ `$> less mytrace.clustered.clusters_info.csv`

- ⌘ % Total duration row, quantifies how 'important' is a cluster.

- ⌘ Cluster numeration ordered by this metric, being Cluster 1 the highest

# Clustering: general tips and issues

⌘ How do I know my clustering 'went well' ?

⌘ "There's no bad clustering as long as you can extract information."

⌘ Compare plot against the clustered trace. Are you happy?

⌘ gnuplot vs trace viewed with clusterID\_window cfg.

⌘ Is my Epsilon too high or too small?

⌘ Lots of noise, or crazy patterns, maybe too small.

⌘ You have only a couple of big clusters and I think should appear more: too high.

⌘ There's no magic rule

⌘ Choose some parameters, run clustering and see results...

⌘ Not good? Refine parameters, repeat until satisfied.

# Clustering: general tips and issues

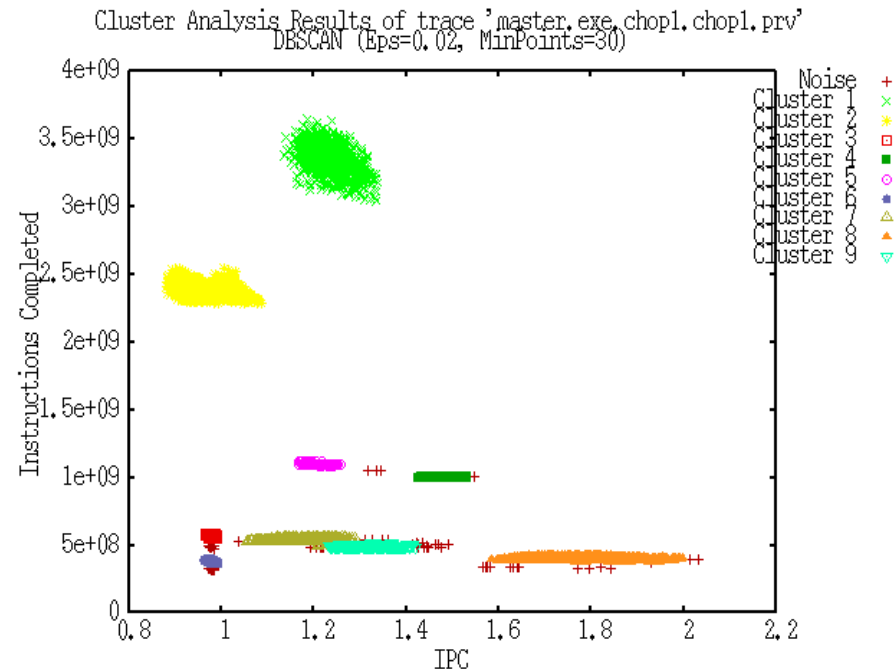
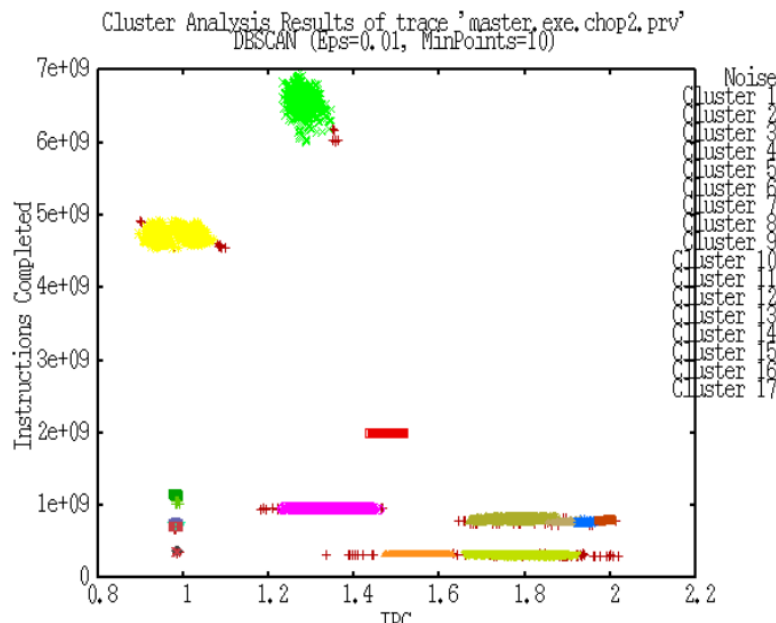
⌘ Important ! If you have several traces scaling number of cores to do tracking later..

⌘ Be careful, the clusters 'move' and maybe they end up out of the min/max ranges.

⌘ Use min/max cluster.xml filters to frame the same regions for each trace. Like a cameraman follows the action.

⌘ The better adjusted, the easier for the tracking tool.

⌘ Look at left image, bottom region, it disappears in the right image, not good!




# Clustering: general tips and issues

« After adjusting the ranges...

« This is better


« This two clustered traces are ready to be used with the tracking tool

Cluster Analysis Results of trace '128.prv'  
DBSCAN (Eps=0.01, MinPoints=10)



0.50000 0.00000

Cluster Analysis Results of trace '256.prv'  
DBSCAN (Eps=0.01, MinPoints=10)





**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

**TRACKING**

# Tracking

- « A friendly way to quantify and visualize the evolution of the clusters among several traces.
- « Tool has two parts
  - « Recognition algorithm 'who is who', based on heuristics.
  - « And visualization GUI
- « Implemented Heuristics
  - « Proximity (default), Callstack (if enabled), density of inst., alignment (experimental)
- « Examples analyzing multiple traces
  - « Scaling number of cores ( 64 – 128 - 256... CPU's)
  - « Testing different microarchitecture features.

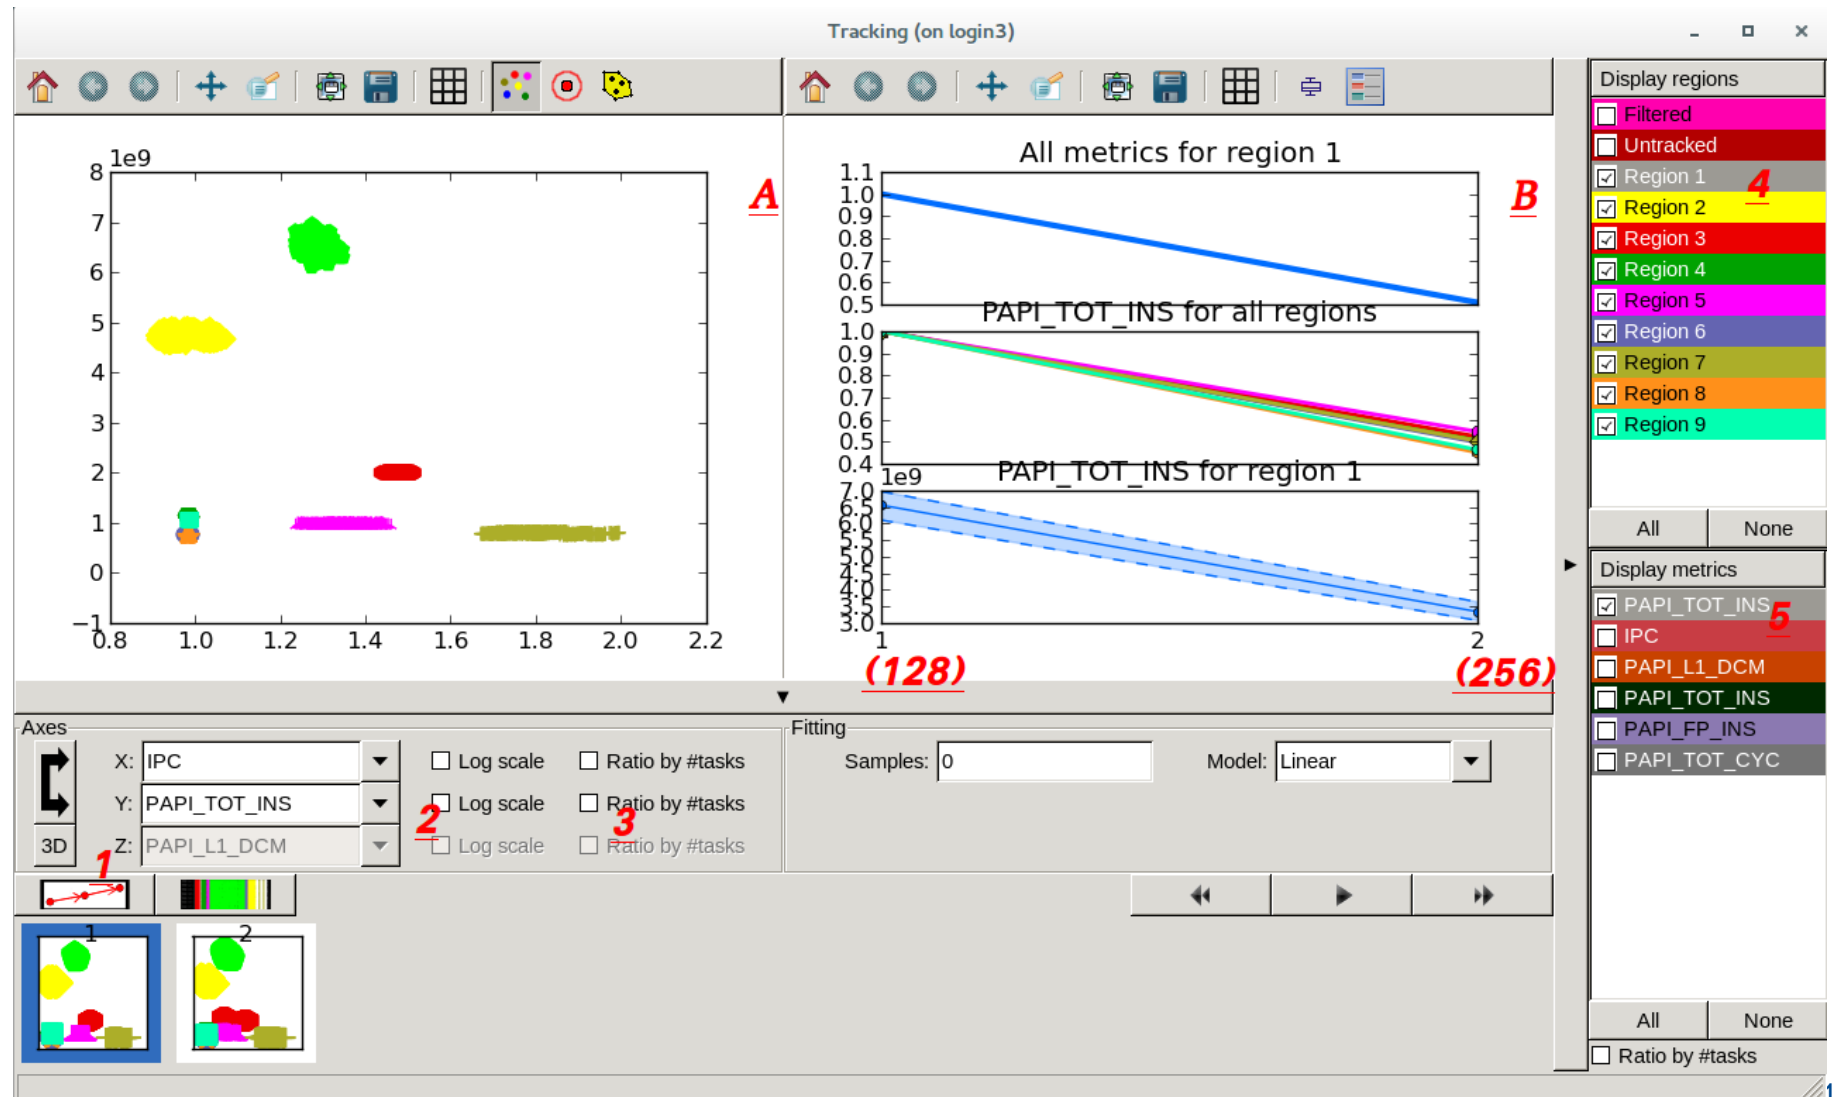


# Tracking – Set up and run

- ⌘ We keep assuming you're connected to MN3
- ⌘ Your starting point is the folder with all the files generated during clustering.
  - ⌘ You applied clustering to 2 traces or more. i.e ( 64.prv - 128.prv - 256.prv )
  - ⌘ Remember: module load BSCTOOLS
- ⌘ Execute
  - ⌘ `tracking -r 128.clustered.prv 256.clustered.prv [ 512... ]`
  - ⌘ Remember to put the traces in logical order
- ⌘ tracking command generates TRACKING.RESULTS.xtrack
  - ⌘ Open with `xtrack TRACKING.RESULTS.xtrack`
  - ⌘ You will also have new clustered traces but with new colors.

# Tracking – GUI

« Comments on next slide



# Tracking – GUI

- ⌘ Window A shows a plot, by default the Instructions / IPC plot for a single trace.
  - ⌘ Push button 1 (dots and arrows). Shows the movement of the clusters along the different traces.
  - ⌘ Now push 2. You can see clearly the movement of the clusters in the lower part of the plot.
  - ⌘ 3 Divides the numbers by number of proceses. Useful to see the impact of code replication
- ⌘ Window B, has another 3 plots, they can be modified using right controls (4,5)
  - ⌘ Note that only shows a straight line because we only used two traces (128,256).
  - ⌘ “All regions” means: all clusters with a check
  - ⌘ “For region 1” means: you clicked/selected Cluster 1 label (it becomes grey)
  - ⌘ In this example, we can see how the Total Instructions were reduced to 50% for cluster one (top plot).
  - ⌘ Hard to explain in slides. Play around with the controls to understand better.

# Thanks!

- « Special thanks to G. Llort @ bsc tools team for teaching me the methodology of use of the tools.
- « Thanks to everyone that gave me feedback after the presentation.
- « Feel free to improve / expand this slides :)



**Barcelona  
Supercomputing  
Center**

*Centro Nacional de Supercomputación*

**Thank you!**

For further information please contact  
[constantino.gomez@bsc.es](mailto:constantino.gomez@bsc.es)